

Module 4: Recurrent Neural Network

To write legible answers you will need to be familiar with both [Markdown](#) and [Latex](#)

Before you turn this problem in, make sure everything runs as expected. First, restart the kernel (in the menubar, select Kernel→Restart) and then run all cells (in the menubar, select Cell→Run All).

Make sure you fill in any place that says "YOUR CODE HERE" or "YOUR ANSWER HERE", as well as your name below:

```
In [ ]: NAME = ""
        STUDENT_ID = ""
```

Some imports we will need:

```
In [ ]: import tensorflow as tf
        from tensorflow.keras.datasets import reuters
        import matplotlib.pyplot as plt
        import numpy as np
        import string
        import textwrap
        from tensorflow.keras.preprocessing.sequence import pad_sequences
        from tensorflow.keras.utils import to_categorical
        from tensorflow.keras.layers import Embedding, Dense, Dropout, Input, LSTM, GRU, Bi
        from tensorflow.keras.models import Model
```

We will be using the [Reuters newswire](#) classification dataset, which has text paired with 46 topics as labels. You can see what these labels represent [here](#).

Let's load the data:

```
In [ ]: (X_train, y_train), (_, _) = reuters.load_data()
```

```
In [ ]: # https://stackoverflow.com/questions/42821330/restore-original-text-from-keras-s-i
        # Needed to encode our own reviews later

        word_dict = reuters.get_word_index()
        word_dict = {k:(v+3) for k,v in word_dict.items()}
        word_dict["<PAD>"] = 0
        word_dict["<START>"] = 1
        word_dict["<UNK>"] = 2
        word_dict["<UNUSED>"] = 3

        vocab_size = len(word_dict.keys())
```

```
In [ ]: # Needed to decode training data into readable text
```

```
inverse_word_dict = {value:key for key,value in word_dict.items()}
```

```
In [ ]: X_train = np.array(X_train)
X_train = pad_sequences(X_train)
```

```
max_sequence_len = X_train[0].shape[0]
print('Padded to longest sequence length: ', max_sequence_len)
```

```
In [ ]: y_train = to_categorical(y_train, 46)
y_train = np.array(y_train)
```

```
In [ ]: print('Number of words in vocabulary: ', vocab_size)
```

```
In [ ]: def encode_text(text, word_dict, maxlen):
    encoded_text = []
    for raw_word in text.split(' '):
        word = raw_word.strip().strip(string.punctuation).lower()
        if word is '' or word is '\n':
            continue
        try:
            encoded_text.append(word_dict[word])
        except KeyError as e:
            # raise KeyError(f'{e} not in word dictionary, text not encoded.')
            continue
    return pad_sequences(np.array(encoded_text).reshape(1,-1), maxlen=maxlen)

def decode_text(encoded_text, inverse_word_dict):
    sentence = []
    for encoded_word in encoded_text:
        if encoded_word == 0:
            continue
        sentence.append(inverse_word_dict[encoded_word])
    w = textwrap.TextWrapper(width=120,break_long_words=False,replace_whitespace=False)
    return '\n'.join(w.wrap(' '.join(sentence)))
```

Let's take a look at an article in our training data:

```
In [ ]: idx = 144
```

```
print(decode_text(X_train[idx], inverse_word_dict), end='\n\n')
```

```
print('Topic: ', y_train[idx])
```

Question 1

Create a model using an RNN layer (LSTM or GRU, unidirectional or bidirectional) to achieve at least 60% validation accuracy in 10 epochs or less:

```
In [ ]: ### YOUR CODE HERE ###
```

Reminder: We have 46 categories. What final activation do we need to use?

Compile your model and display the summary:

```
In [ ]: loss = ### YOUR CODE HERE ###

opt = ### YOUR CODE HERE ###

metrics = ### YOUR CODE HERE ###

reuters_model.compile(loss=loss,
                      optimizer=opt,
                      metrics=metrics)

reuters_model.summary()
```

Train your model:

```
In [ ]: batchsize = ### YOUR CODE HERE ###

history = reuters_model.fit(X_train, y_train, batch_size=batchsize, epochs=10, vali
```

Plot the training and validation losses and accuracies:

```
In [ ]: ### YOUR CODE HERE ###
```

Question 2

We've seen both LSTM and GRU cells as building blocks for RNNs.

Here is a reminder of each of their corresponding architectures:



- i) What are the major differences between each?
- ii) What are the major advantages of each?
- iii) What are the major disadvantages of each?

YOUR ANSWERS HERE